



TECHNISCHE
UNIVERSITÄT
WIEN

Exercise 3: Sockets

191.002 Operating Systems VU
2024W

Florian Mihola, Axel Brunnbauer, Luca Di Stefano, Peter Puschner
2024-11-19

Overview

Inter-process communication

Considered so far...

Exchanging data between processes on the same system

- Explicit synchronization between unrelated processes
 - Shared Memory
 - Semaphores
- Implicit synchronization between related processes
 - Blocking read- and write operations
 - Non-related processes via sockets
 - Related processes via unnamed pipes

Today...

Exchanging data via sockets - either on the same system or over a network

- Implicit synchronization between unrelated processes

Byte Order or Endianness

- Sequential ordering of bytes in memory

```
1  int i = 0x12345678; // 8 hex digits = 4 bytes
```

- Little endian: little end first = least significant byte first

Byte address	$\&i$	$\&i + 1$	$\&i + 2$	$\&i + 3$
Byte content	0x78	0x56	0x34	0x12

- Big endian: big end first = most significant byte first

Byte address	$\&i$	$\&i + 1$	$\&i + 2$	$\&i + 3$
Byte content	0x12	0x34	0x56	0x78

- Byte order in memory depends on processor architecture (x86 is little endian)
- When writing multiple bytes, program must take care of byte order
- Network byte order is big endian

Byte Order or Endianness

Write bytes explicitly in little endian order:

```
1  int i = 0x12345678;
2  uint8_t buf[sizeof(int)];
3  int pos;
4  for (pos = 0; pos < sizeof(int); pos++)
5      buf[pos] = i >> 8 * pos;
6
7  fwrite(buf, sizeof(int), 1, out);
```

Read bytes explicitly in little endian order:

```
1  uint8_t buf[sizeof(int)];
2  fread(buf, sizeof(int), 1, in);
3  int i = 0;
4  int pos;
5  for (pos = 0; pos < sizeof(int); pos++)
6      i |= (int)buf[pos] << 8 * pos;
7  // i == 0x12345678
```

Byte Order or Endianness

`uint32_t htonl(uint32_t netlong)`

- Convert a 32-bit from host byte order to network byte order

`uint32_t ntohl(uint32_t netlong)`

- Convert a 32-bit integer from network byte order to host byte order

Sockets

- What is a socket?
 - Method for interprocess communication (IPC)
 - Either on a single host or between different hosts in a network (or via internet)
- Common scenario: communication between a client and a server
- Sockets are handled like files
 - Each socket gets a file descriptor
 - Reading and writing to the associated file descriptor

Socket API

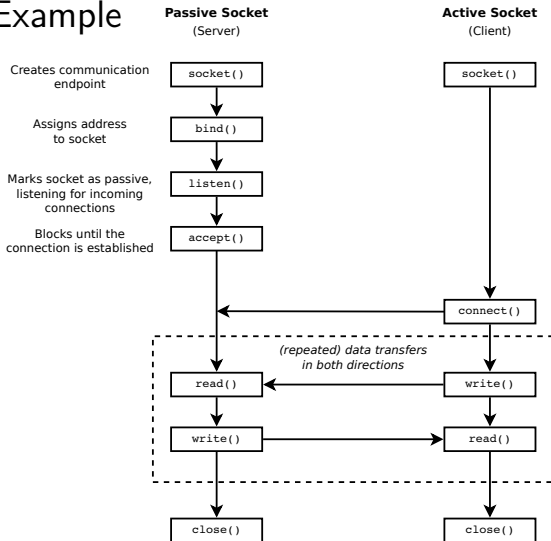
- Sockets are an interface to the transport layer of a communication protocol
 - Direct communication between client and server: no need to know the network layout
 - Sockets do not implement application protocols (HTTP, FTP, ...)
- Connection-oriented, bidirectional and reliable communication channel
- The connection is established between two endpoints
 - Endpoint on server side: Server IP + known port number
 - Endpoint on client side: Client IP + unused port number

Sockets

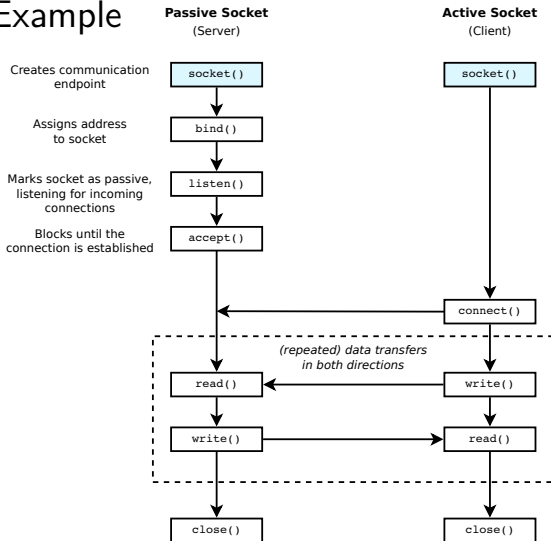
Address families and socket types

- Address family (network layer)
 - Internet Protocol, version 4 (IPv4)
AF_INET → man 7 ip
 - Internet Protocol, version 6 (IPv6)
AF_INET6 (IPv6) → man 7 ipv6
 - Unix Domain Sockets (local IPC)
AF_UNIX → man 7 unix
- Socket type
 - Connection-oriented sockets (stream based)
 - SOCK_STREAM, default for IP is TCP
 - Connection is identified by two endpoints
 - Connection-less sockets (datagram/message based)
 - SOCK_DGRAM, default for IP is UDP

Client-Server Example



Client-Server Example



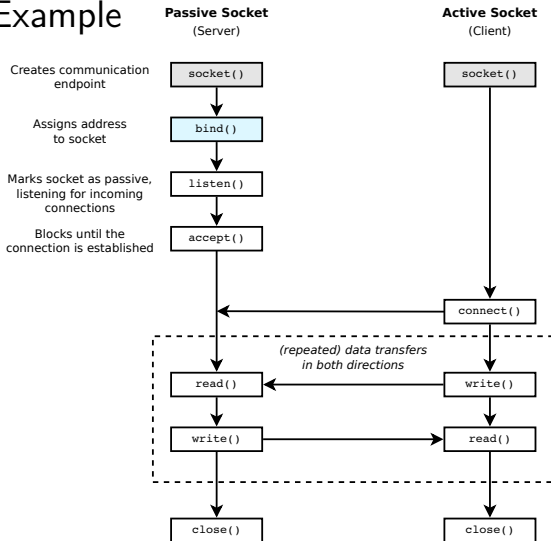
System Call: socket()

`int socket(int family, int type, int protocol)`

- Creates a communication endpoint (socket)
 - `family` address family
 - `type` socket type
 - `protocol` communication protocol to be used
 - address family + type usually implies protocol
 - 0 for default-protocol
- Return value: `File descriptor` of the newly created socket or -1 on failure (→ `errno`)

```
1 int sockfd = socket(AF_INET, SOCK_STREAM, 0);
2
3 if (sockfd < 0)
4     // error, check errno for details
```

Client-Server Example



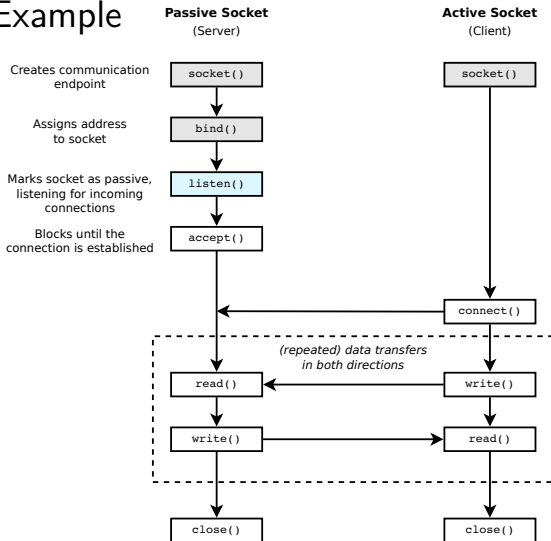
System Call: bind()

```
int bind(int socket, struct sockaddr *address,  
         socklen_t addr_len)
```

- Assigns the specified address to a socket
 - `socket` file descriptor of the socket
 - `address` data structure with the desired address
 - `addr_len` size of the address data structure
- Return value: 0 on success, -1 on failure (\rightarrow errno)

```
1  struct sockaddr_in *sa;  
2  ...  
3  
4  if (bind(sockfd, sa, sizeof(struct sockaddr_in)) < 0)  
5      // error
```

Client-Server Example



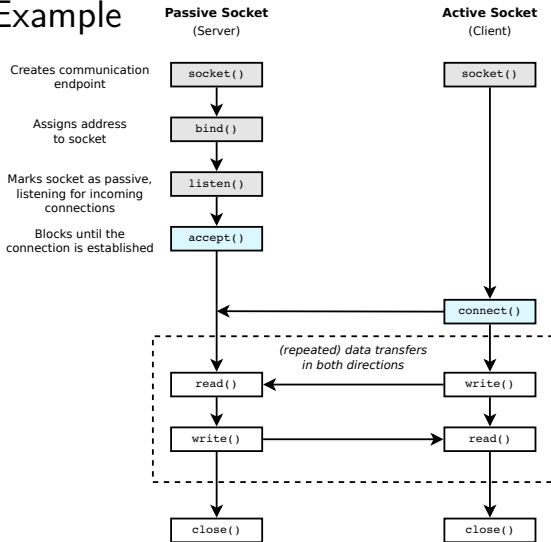
System Call: listen()

`int listen(int socket, int backlog)`

- Listen for connections on a socket (= mark it as passive)
- For connection-oriented protocols only
 - `socket` socket file descriptor
 - `backlog` number of connection requests, which are managed in a queue by the OS, until the server accepts them
- Return value: 0 on success, -1 on failure (→ `errno`)

```
1  if (listen(sockfd, 1) < 0)
2      // error
```

Client-Server Example



System Call: accept()

```
int accept(int socket, struct sockaddr *address,  
           socklen_t *addr_len)
```

- Accept a new connection on a socket (passive, server)
 - socket** socket file descriptor
 - address** pointer to a sockaddr structure where the address of the connecting socket is returned (actual type depends on protocol, e.g. `sockaddr_in`), `NULL` possible
 - addr_len** pointer to the size of the structure in address
- Blocks if there is no pending request
- Returns a new socket (file descriptor) for the first pending connection or -1 on error (\rightarrow `errno`)

```
1 int connfd = accept(sockfd, NULL, NULL);  
2 if (connfd < 0)  
3     // error
```

System Call: connect()

```
int connect(int socket, const struct sockaddr *address,  
            socklen_t addr_len)
```

- Initiate a connection (active, client)
 - `socket` socket file descriptor
 - `address` address of the server (destination)
 - `addr_len` size of the address structure
- Returns after the connection has been established
- The operating system of the client selects an arbitrary, unused port

```
1 struct sockaddr_in server_addr;  
2 ...  
3  
4 if (connect(sockfd, &server_addr, sizeof(server_addr)) < 0)  
5     // error
```

getaddrinfo(3)

int **getaddrinfo**(const char *node, const char *service,
const struct addrinfo *hints, struct addrinfo **res)

- Create a suitable socket address with getaddrinfo(3)
 - node** Hostname (e.g. "localhost", "173.194.44.232", "google.com") or NULL (for usage with bind())
 - service** port no. or name of service (e.g. "80", "http")
 - hints** Selection criteria
 - res** Destination address for the resulting addrinfo structure (filled by getaddrinfo)
- Returns 0 on success or an error code (**no use of errno!**)
- See also **gai_strerror(3)** and **freeaddrinfo(3)**

```
1 struct addrinfo hints, *ai;  memset(&hints, 0, sizeof(hints));
2 hints.ai_family = AF_INET;
3 hints.ai_socktype = SOCK_STREAM;
```

getaddrinfo(3)

`int getaddrinfo(const char *node, const char *service,
const struct addrinfo *hints, struct addrinfo **res)`

- Create a suitable socket address with `getaddrinfo(3)`
 - `node` Hostname (e.g. "localhost", "173.194.44.232", "google.com") or NULL (for usage with `bind()`)
 - `service` port no. or name of service (e.g. "80", "http")
 - `hints` Selection criteria
 - `res` Destination address for the resulting `addrinfo` structure (filled by `getaddrinfo`)
- Returns 0 on success or an error code (**no use of `errno`!**)
- See also `gai_strerror(3)` and `freeaddrinfo(3)`

```
1 struct addrinfo hints, *ai;  memset(&hints, 0, sizeof(hints));
2 hints.ai_family = AF_INET;
3 hints.ai_socktype = SOCK_STREAM;
4 int res = getaddrinfo("localhost", "1280", &hints, &ai);
5 if (res != 0) { fprintf(stderr, "getaddrinfo: %s\n", gai_strerror(res)); }
```

Example: getaddrinfo()

Client

```
1  struct addrinfo hints, *ai;
2  memset(&hints, 0, sizeof hints);
3  hints.ai_family = AF_INET;
4  hints.ai_socktype = SOCK_STREAM;
5
6  int res = getaddrinfo("localhost", "1280", &hints, &ai);
7  if (res != 0) {
8      // error
9  }
10 int sockfd = socket(ai->ai_family, ai->ai_socktype,
11                    ai->ai_protocol);
12 if (sockfd < 0) {
13     // error
14 }
15 if (connect(sockfd, ai->ai_addr, ai->ai_addrlen) < 0) {
16     // error
17 }
18 freeaddrinfo(ai);
```

Example: getaddrinfo()

Server

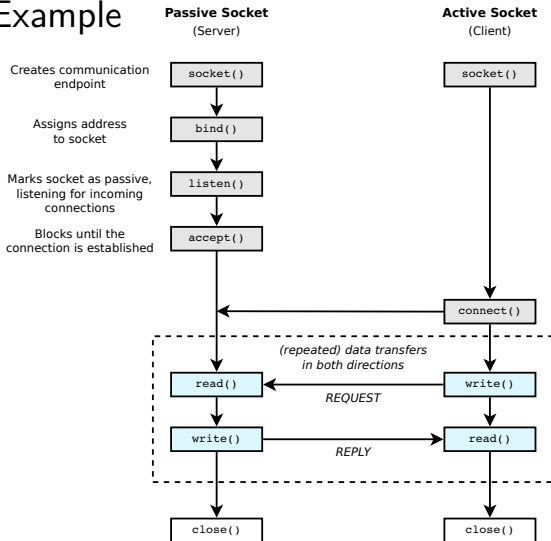
```
1  struct addrinfo hints, *ai;
2  memset(&hints, 0, sizeof hints);
3  hints.ai_family = AF_INET;
4  hints.ai_socktype = SOCK_STREAM;
5  hints.ai_flags = AI_PASSIVE;                                // <--
6  int res = getaddrinfo( NULL, "1280", &hints, &ai);         // <--
7  if (res != 0) {
8      // error
9  }
10 int sockfd = socket(ai->ai_family, ai->ai_socktype,
11                    ai->ai_protocol);
12 if (sockfd < 0) {
13     // error
14 }
15 if ( bind(sockfd, ai->ai_addr, ai->ai_addrlen) < 0) {        // <--
16     // error
17 }
18 freeaddrinfo(ai);
```

gethostbyname(3)

getaddrinfo replaces the obsolete function `gethostbyname`

- `gethostbyname` does not support IP version 6 and is obsolete
- Most of the C socket examples that can be found online still use the old `gethostbyname`
- **You must not use `gethostbyname` and related functions (i.e. `gethostbyaddr`, `gethostbyname2`, `gethostent_r`, `gethostbyaddr_r`, `gethostbyname_r`, `gethostbyname2_r`, ...) during the exercises or the exams!**

Client-Server Example



Send and Receive

write(2) and read(2)

- After the connection has been established, the **file descriptor** of the socket is used to read and write data
- Use **read** and **write** the same way as with files

```
1  char buf[80];
2  int pos, cnt;
3
4  for (pos = 0; pos < sizeof(buf); ) {
5      cnt = read(sockfd, buf + pos, sizeof(buf) - pos);
6      if (cnt < 0) {
7          if (errno != EINTR) { /* other error than EINTR */ }
8      } else
9          pos += cnt;
10 }
```

- You can also use the Stream I/O with **fdopen()** (don't forget about buffering, use **fflush()** to send the data!)

Send/Receive using Stream I/O - Example (no error handling)

```
1 struct addrinfo hints, *ai;
2 memset(&hints, 0, sizeof(hints));
3 hints.ai_family = AF_INET;
4 hints.ai_socktype = SOCK_STREAM;
5 getaddrinfo("neverssl.com", "http", &hints, &ai);
6
7 int sockfd = socket(ai->ai_family, ai->ai_socktype,
8                    ai->ai_protocol);
9 connect(sockfd, ai->ai_addr, ai->ai_addrlen);
```

Send/Receive using Stream I/O - Example (no error handling)

```
1 struct addrinfo hints, *ai;
2 memset(&hints, 0, sizeof(hints));
3 hints.ai_family = AF_INET;
4 hints.ai_socktype = SOCK_STREAM;
5 getaddrinfo("neverssl.com", "http", &hints, &ai);
6
7 int sockfd = socket(ai->ai_family, ai->ai_socktype,
8                     ai->ai_protocol);
9 connect(sockfd, ai->ai_addr, ai->ai_addrlen);
10 // send and receive data over connection
11 FILE *sockfile = fdopen(sockfd, "r+");
12
13 fputs("GET / HTTP/1.1\r\nHost: neverssl.com\r\n\r\n",
14       sockfile);
15 fflush(sockfile); // send all buffered data
16
17 char buf[1024];
18 while (fgets(buf, sizeof(buf), sockfile) != NULL)
19     fputs(buf, stdout);
```

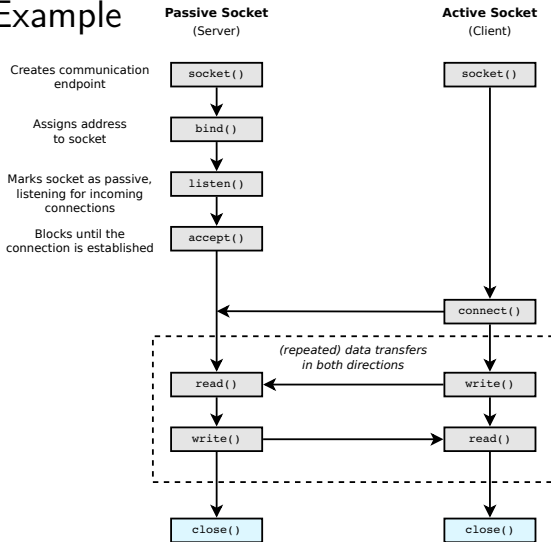
Send and Receive

send(2) and recv(2)

```
int send(int socket, const void *msg, size_t msg_len, int flags)
int recv(int socket, void* buf, size_t buf_len, int flags)
```

- Specializations of **write** und **read** for sockets
- Return value and first three arguments same as for **write** und **read**
- Additional argument: **flags**
 - MSG_DONTWAIT – Non-blocking send/receive
 - MSG_WAITALL – Block until all data was received (exceptions: error, signal received)

Client-Server Example



Socket Options

```
int setsockopt(int socket, int level, int option_name,  
               const void *option_value, socklen_t option_len)
```

- Set options on a socket (see man page for full list: `setsockopt(2)`, `socket(7)`, `ip(7)`)
- Useful to avoid the error “Address already in use” (`EADDRINUSE`) with `bind` upon restarting your server program (otherwise the port remains unusable for approximately 1 min after the server was terminated)

```
1 int optval = 1;  
2 setsockopt(serverfd, SOL_SOCKET, SO_REUSEADDR, &optval,  
3           sizeof optval);
```

Exercise 3

Client and server for HTTP

- 3A: Client
- 3B: Server
- IPC via stream-oriented sockets
- Implement a subset of the HTTP (HyperText Transfer Protocol), used for requesting websites
- Your server can serve files to a web browser (e.g. Firefox)
- Your client can request files from webserver (unfortunately most webserver require HTTPS)
 - `http://pan.vmars.tuwien.ac.at/osue/`
 - `http://neverssl.com/`
 - `http://www.nonhttps.com/`

Material

- The GNU C Library Reference Manual,
Ch. 12 (Stream I/O), Ch. 16 (Sockets)
http://www.gnu.org/software/libc/manual/html_node/
- Beej's Guide to Network Programming
<http://beej.us/guide/bgnet/>