



TECHNISCHE
UNIVERSITÄT
WIEN

Exercise 2: Processes & Pipes

191.002 Operating Systems VU
2024W

Florian Mihola, Axel Brunnbauer, Luca Di Stefano, Peter Puschner
2024-11-12

Processes

Why should we create processes?

- Divide up a task
 - Simpler application design
 - Greater concurrency


Example

A server listens to client requests. The server process starts a new process to handle each request and continues to listen for further connections. The server can handle several client requests simultaneously.

Process vs. Thread

`fork(2)` vs. `pthread(7)`


Process 0




Process vs. Thread

fork(2) vs. pthreads(7)

Process 0




Process 1



Process vs. Thread

`fork(2)` vs. `pthread(7)`


Process 0




Process vs. Thread

`fork(2)` vs. `pthread(7)`

Process 0



Process 1



Process Hierarchy

- Every process has a parent process
- Exception: init process (init, systemd)
- Every process has a unique ID (pid_t)
- Show process hierarchy: `ps tree(1)`

```
systemd-+-ModemManager---2*[{ModemManager}]
        |-NetworkManager--dhclient
        |               '-2*[{NetworkManager}]
        |-abrt-dbus---{abrt-dbus}
        |-2*[abrt-watch-log]
        |-abrt-d
        |-acpid
        |-agetty
        |-alsactl
        |-atd
        |-auditd-+-audispd-+-sedispatch
        |               |               '-{audispd}
        |               '-{auditd}
        |-automount---7*[{automount}]
        |-avahi-daemon---avahi-daemon
        |-chronyd
        |-colord---2*[{colord}]
        |-crond
        |-cupsd
        |-dbus-daemon
        |-dnsmasq---dnsmasq
        |-firewalld---{firewalld}
        .
        .
```

Properties of a Process in Linux

State Running, waiting, ...

Scheduling Priority, CPU time, ...

Identification PID, owner, group, ...


Memory Management Pointer to MMU information

Signals Mask, pending

Process Relations Parents, siblings

- Show process info: `cat /proc/<pid>/status`
- See `struct task_struct` in `sched.h`


virtual memory
of a process



Interface

fork / exec / exit / wait


- **fork(2)** – creates a process (copies the process image)
- **exec(3)** – loads a program (replaces the process image of a process with a new one)
- **exit(3)** – exits a process
- **wait(2)** – awaits the exit of child processes



Process Creation

fork

- Creates a new process
- New process is an identical copy of the calling process – except PID, pending signals, ...
- Calling process is the **parent** of the created process, the **child** – processes are **related**
- Both processes run parallel and execute the same program (from the **fork** call on)



Process Creation

fork


- Create the process

```
1  #include <unistd.h>
2
3  pid_t fork(void);
```


- Distinguish between parent and child via return value of `fork`
 - 1 On error
 - 0 In the child process
 - >0 In the parent process

Process Creation

Before `fork()`



After `fork()`



Process Creation

Example

```
1  pid_t pid = fork();
2
3  switch (pid) {
4      case -1:
5          fprintf(stderr, "Cannot fork!\n");
6          exit(EXIT_FAILURE);
7
8      case 0:
9          [...] // child tasks
10         break;
11     default:
12         [...] // parent tasks
13         break;
14 }
```

Process Creation

Child

Child inherits from parent:

- Opened files (common access!)
- File buffers
- Signal handling
- **Current** values of variables


But:

- Variables are local to process (no influence)
- Signal handling can be re-configured
- Communication (IPC) via pipes, sockets, shared memory, ...

Program Execution

exec

- Load a new program into a process's memory
- Executes **another** program
- In the **same** process (PID remains the same)



Program Execution

exec Family¹

```
1  int execl(const char *path, const char *arg, ...);
2  int execlp(const char *file, const char *arg, ...);
3
4  int execlx(const char *path, const char *arg, ...,
5             char *const envp[]);
6
7  int execv(const char *path, char *const argv[]);
8  int execvp(const char *file, char *const argv[]);
9
10 int fexecve(int fd, char *const argv[], char *const envp[]);
```

¹Frontend of execve(2)

Program Execution

exec Family

- `exec``l``□` – variable number of arguments
- `exec``v``□` – arguments via array
- `exec``□``p` – searching the environment variable `$PATH` for the program specified
- `exec``□``e` – environment² can be changed
- `fexecve` – accepts file descriptor (instead of path)

Note Argument Passing!

- 1st argument is the program's name (`argv[0]`)!
- Last argument must be a **NULL** pointer!

²FYI: `environ(7)`

Program Execution

Example: `execv()`, `execvp()`

```
1  #include <unistd.h>
2
3
4  char *cmd[] = { "ls", "-l", (char *) 0 };
5
6  execv("/bin/ls", cmd);
7  // or:
8  // execvp("ls", cmd);
9
10 fprintf(stderr, "Cannot exec!\n");
11 exit(EXIT_FAILURE);
```

Program Execution

Example: `execl()`, `execlp()`

```
1  #include <unistd.h>
2
3  execl("/bin/ls", "ls", "-l", NULL);
4  // or:
5  // execlp("ls", "ls", "-l", NULL);
6  fprintf(stderr, "Cannot exec!\n");
7  exit(EXIT_FAILURE);
```


Attention - this will not work:

```
1  execl("/bin/ls", "ls -l", NULL);
2
3  int a = 1;
4  execl("myprog", "myprog", "-a", a, NULL);
5      // e.g., use a char-buffer and snprintf(3)
```

Process Termination

exit

- Terminates a process (normally)
- Termination status can be read by parents
- Actions performed by `exit()`
 - Flush and close stdio stream buffers
 - Close all open files
 - Delete temporary files (created by `tmpfile(3)`)
 - Call exit handlers (`atexit(3)`)



Process Termination

exit

- Terminate a process normally


```
1 #include <stdlib.h>
2
3 void exit(int status);
```

- Status: 8 bit (0-255)
- By convention
 - `exit(EXIT_SUCCESS)` – process completed successfully
 - `exit(EXIT_FAILURE)` – error occurred
- More return values
 - BSD: `sysexits.h`
 - <http://tldp.org/LDP/abs/html/exitcodes.html>

Waiting on a Child Process

wait

- Wait until a child process terminates
- Returns the PID and status of the terminated child



Waiting on a Child Process

wait

- Wait for a child to terminate

```
1 #include <sys/wait.h>
2
3 pid_t wait(int *status);
```

- `wait()` blocks³ until a child terminates or on error
- Return value
 - PID of the terminated child
 - -1 on error (→ `errno`, e.g., `ECHILD`)
- Status includes exit value and signal information
 - `WIFEXITED(status)`, `WEXITSTATUS(status)`
 - `WIFSIGNALED(status)`, `WTERMSIG(status)`
 - See `wait(2)`

³≠ busy waiting

Waiting on a Child Process

Zombies and Orphans

- UNIX: Terminated processes remain in the process table
- No more space in process table → no new process can be started!
- After `wait()` the child process is removed from the process table

Zombie Child terminates, but parent didn't call `wait` yet

- State of the child is set to "zombie"
- Child remains in process table until parent calls `wait`

Orphan Parent terminates before child

- Child becomes an **orphan** and is inherited to the `init` process
- When an orphan terminates, the `init` process removes the entry in the process table

Waiting on a Child Process

Example

```
1  #include <sys/wait.h>
2
3  int status;
4  pid_t child_pid, pid;
5  ...
6  while ((pid = wait(&status)) != child_pid) {
7      if (errno == EINTR) continue;
8      if (pid == -1) {
9          fprintf(stderr, "Cannot wait!\n");
10         exit(EXIT_FAILURE);
11     }
12 }
13
14 if (WEXITSTATUS(status) == EXIT_SUCCESS) {
15     ...
```

Waiting on a Child Process

waitpid

- Wait on a **specific** child process

```
1 #include <sys/wait.h>
2
3 pid_t waitpid(pid_t pid, int *status, int options);
```

- Examples

```
1 waitpid(cid, &status, 0);
2     // waits on a child process with PID 'cid'
3
4 waitpid(-1, &status, 0);
5     // equivalent to wait
6
7 waitpid(-1, &status, WNOHANG);
8     // does not block
```

Notification

on Termination of a Child

If parent should not block

- Synchronous
 - `waitpid(-1, &status, WNOHANG)`
 - Returns exit status when a child terminates
 - Repeating calls → polling
- Asynchronous
 - Signal `SIGCHLD` is sent to the parent process whenever one of its child processes terminates
 - Catch by installing a signal handler (`sigaction`)
 - Call `wait` in the signal handler

Pitfalls

```
1  int main(int argc, char **argv)
2  {
3      fprintf(stdout, "Hello");
4
5      (void) fork();
6      return 0;
7  }
```

Output: "HelloHello"

Why?

Pitfalls

```
1  int main(int argc, char **argv)
2  {
3      fprintf(stdout, "Hello");
4      fflush(stdout);
5      (void) fork();
6      return 0;
7  }
```

Output: "Hello"

→ for all opened streams

Debugging

`gdb`

- Before `fork` is executed:
`set follow-fork-mode [child|parent]`

Example

```
1  $ gdb -tui ./forktest
2  (gdb) break main
3  (gdb) set follow-fork-mode child
4  (gdb) run
5  (gdb) next
6  (gdb) :
7  (gdb) continue
8  (gdb) quit
```

Inter-Process Communication

Recall

So far:

- Signals (e.g., to synchronise between parent and child)
→ see [Development in C I](#)

New:

- Pipes

Other lectures:

- Shared Memory
- Sockets


Pipes

(Unnamed) Pipe

= unidirectional data channel

= enables communication between **related** processes

- Example: `$ ls | wc -l`



- Access to read and write end of the pipe via file descriptors
- Pipe is an unidirectional byte stream
- Buffered
- Implicit synchronisation

Pipes

Create

- Create a pipe

```
1 #include <unistd.h>
2
3 int pipe(int pipefd[2]);
```


- File descriptors of read and write end are returned in specified integer array `pipefd`
 - `pipefd[0]` – read end
 - `pipefd[1]` – write end
- Close unused ends
- Use read/write end via stream-IO (`fdopen`, etc.)
- A child process inherits the pipe → common access

Unnamed Pipes

Illustration

```
pipe;
```

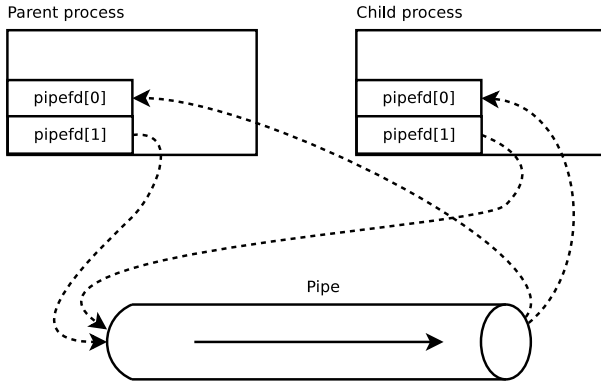
Parent process



Unnamed Pipes

Illustration


```
pipe;  
fork;
```



Unnamed Pipes

Illustration

```
pipe;  
fork;  
close unused ends;
```



Unnamed Pipes

Implicit Synchronisation

- `read` blocks on empty pipe
- `write` blocks on full pipe
- `read` indicates **end-of-file** if all write ends are closed (return value 0)
- `write` creates signal **SIGPIPE** if all read ends are closed (if signal ignored/handled: `write` fails with `errno EPIPE`)

Therefore...

... close unused ends, to get this behaviour (end-of-file and **SIGPIPE/EPIPE**).

Besides, the kernel removes pipes with all ends closed.

Unnamed Pipes

What about named pipes?

- Unnamed pipes
 - `|`
 - `pipe(2)`
- Named pipes
 - `mkfifo(1)`, `mknod(2)`
 - Usage similar to files.
 - (Will not be dealt with any further throughout this course.)

Redirection of stdin/stdout

Why?

- Main application: pipes
- Example: shell redirection of `stdin` and `stdout`

Scenario:

- A process may be forked or not
→ uses standard IO
- A parent process forks and executes another program
- Parent usually wants to use the child's output
→ redirect `stdin` (file descriptor 0, `STDIN_FILENO`) and/or `stdout` (file descriptor 1, `STDOUT_FILENO`) in new process

Redirection of stdin/stdout

Approach

- Close file descriptors for standard I/O (`stdin`, `stdout`)
- Duplicate opened file descriptor (e.g., a pipe's end) to the closed one

```
1 #include <unistd.h>
2
3 int dup(int oldfd);
4 int dup2(int oldfd, int newfd);
```

- Close duplicated file descriptor

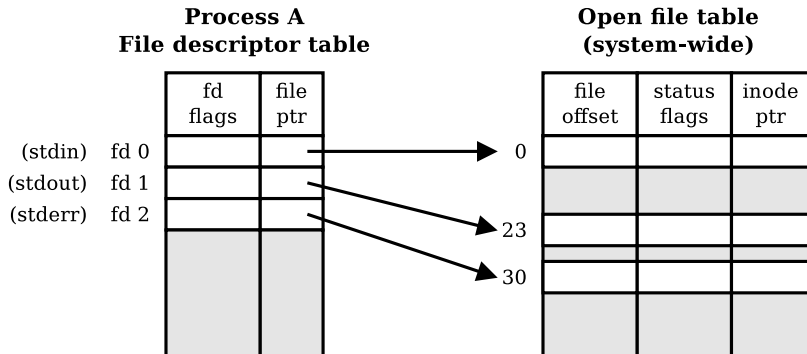
Redirection of stdin/stdout

dup / dup2

- `dup(oldfd)` duplicates file descriptor `oldfd`
 - New file descriptor uses smallest unused ID
= entry in [file descriptor table](#)
 - Duplicated file descriptor points to the [same](#) open file description (equal file offset, status flags) → see `open(2)`
- `dup2(oldfd, newfd)` duplicates `oldfd`
 - New file descriptor uses ID `newfd`
 - (Implicitly) closes the file descriptor `newfd` (if necessary)
 - `newfd` points to the [same](#) open file description like `oldfd`

Redirection of stdin/stdout

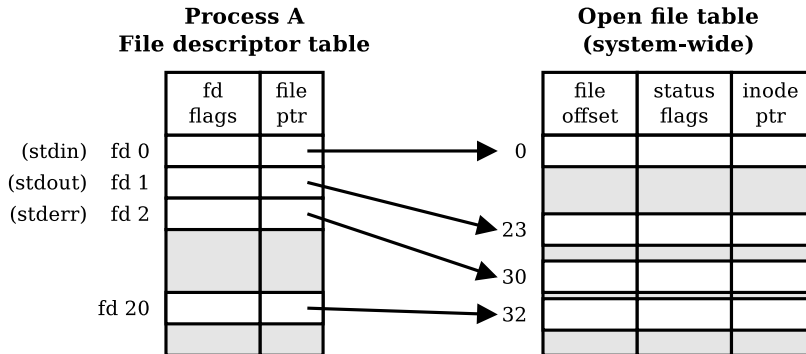
Example: redirect stdout to opened file



Redirection of stdin/stdout

Example: redirect stdout to opened file

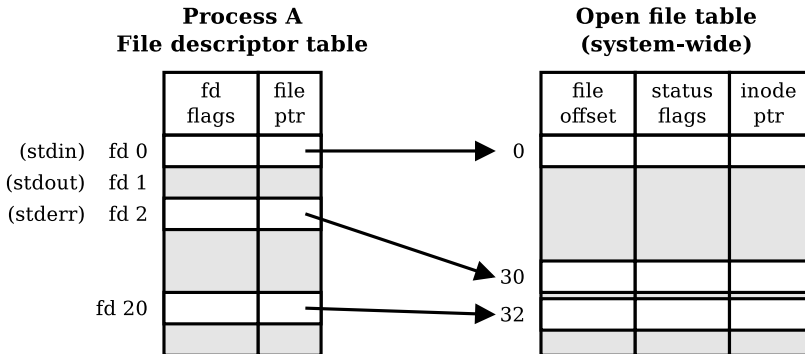
open file;



Redirection of stdin/stdout

Example: redirect stdout to opened file


open file; close stdout;



Redirection of stdin/stdout

Example: redirect stdout to opened file

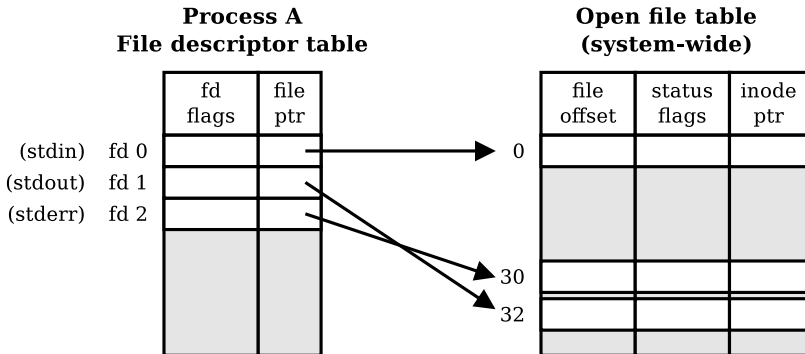
open file; close stdout; dup;



Redirection of stdin/stdout

Example: redirect stdout to opened file

open file; close stdout; dup; close file;



Redirection of stdin/stdout

Example: redirect stdout to log.txt

```
1  #include <fcntl.h>
2  #include <sys/types.h>
3  #include <unistd.h>
4
5  int fd;
6
7  // TODO error handling!
8
9  fd = open("log.txt", O_WRONLY | O_CREAT);
10
11  dup2(fd,          // old descriptor
12       STDOUT_FILENO); // new descriptor
13
14  close(fd);
15
16  execlp("ls", "ls", NULL);
```

Redirection of stdin/stdout

Example: redirect stdin to pipe

```
1  // TODO error handling!
2  int pipefd[2];
3  pipe(pipefd);           // create pipe
4
5  pid_t pid = fork();
6  switch(pid) {
7      [...]
8      case 0: // child counting lines from parent
9          close(pipefd[1]); // close unused write end
10
11          dup2(pipefd[0],    // old descriptor - read end
12              STDIN_FILENO); // new descriptor
13          close(pipefd[0]);
14
15          execlp("wc", "wc", "-l", NULL);
16          // should not reach this line
```


Pitfalls

- Pipes are **unidirectional**
- Bidirectional: two pipes, but ...
 - Erroneous synchronisation (deadlock, e.g., both processes **read** from empty pipe)
- Synchronisation & Buffer
 - Use `fflush()`
 - Configure buffer (`setbuf(3)`, `setvbuf(3)`)

Pipe & Fork

- “Yes, a programmer’s strength flows from fork(). But beware of the dark side.”
- fork is no ordinary function

```
1  int pipefd[2];
2  pipe(pipefd);           // 1. create pipe
3  pid_t pid = fork();    // 2. fork
4  // two processes
5  // one pipe
```

```
1  pid_t pid = fork();    // 1. fork
2  // two processes
3  int pipefd[2];
4  pipe(pipefd);          // 2. create pipeS!
5  // two pipes!
```

Tips for the Exercise

- Try to parallelize the functionality of your program (as much as possible)

Example

DO NOT: The parent first reads all input from a file to an array. It then sends the data within one burst to the child. The child processes the data and outputs the result.

INSTEAD DO: The parent reads line-by-line from a file. Each line is sent to the client immediately. Reading and processing of the lines happens in parallel.

Tips for the Exercise

- Communicate over pipes (do not exploit inherited memory areas)

Example

DO NOT: The parent reads a file and saves its content into an array and forks a child. The child processes the data from the array.

INSTEAD DO: The parent communicates the data from the file over a pipe.

- However, you may pass options/flags/settings to the child (process). For example, use inherited variable `argv` to set arguments when using `exec`.

Material

- Michael Kerrisk: A Linux and UNIX System Programming Handbook, No Starch Press, 2010.
- man pages: `fork(2)`, `exec(3)`, `execve(2)`, `exit(3)`, `wait(3)`, `pipe(2)`, `dup(2)`
- gdb - Debugging Forks:
<https://sourceware.org/gdb/onlinedocs/gdb/Forks.html>